

# Problem 1

(a)

$$R_{tr}(\hat{(\beta)}) = \frac{1}{N} \sum_{i=1}^n (y_i - \beta^T x_i)^2 < \frac{1}{N-p} \sum_{i=1}^n (y_i - \beta^T x_i)^2 = \text{MSE}$$

$$\Rightarrow E(R_{tr}) < E(\text{MSE}) = \sigma^2$$

(b)

$$E(R_{te}(\hat{(\beta)})) = \frac{1}{M} E \sum_{i=1}^M (\tilde{y}_i - \hat{\beta}^T \tilde{x}_i)^2$$

$$= \frac{1}{M} E \left[ \sum_{i=1}^M \left( (\tilde{y}_i - \hat{\beta}^T \tilde{x}_i) + (\beta^T \tilde{x}_i - \beta^T \tilde{x}_i) \right)^2 \right]$$

$$= \frac{1}{M} E \left[ \sum_{i=1}^M (\tilde{y}_i - \beta^T \tilde{x}_i)^2 + \sum_{i=1}^M (\hat{\beta}^T \tilde{x}_i - \beta^T \tilde{x}_i)^2 \right]$$

$$= \frac{1}{M} E \sum_{i=1}^M (\tilde{y}_i - \beta^T \tilde{x}_i)^2 + \frac{1}{M} E \sum_{i=1}^M (\hat{\beta}^T \tilde{x}_i - \beta^T \tilde{x}_i)^2$$

$$= \sigma^2 + (\text{something squared}) > \sigma^2$$

- (c) On average, the training set error will be *less* than the testing set error. This has many implications in both prediction and inference.
- (d) The attached code generates random noise data sets, calculates training and test set risks, then makes a histogram based on 100 simulations. The blue line is  $\sigma^2 = 1$ , and the red line is the average risk for training and test set.

```

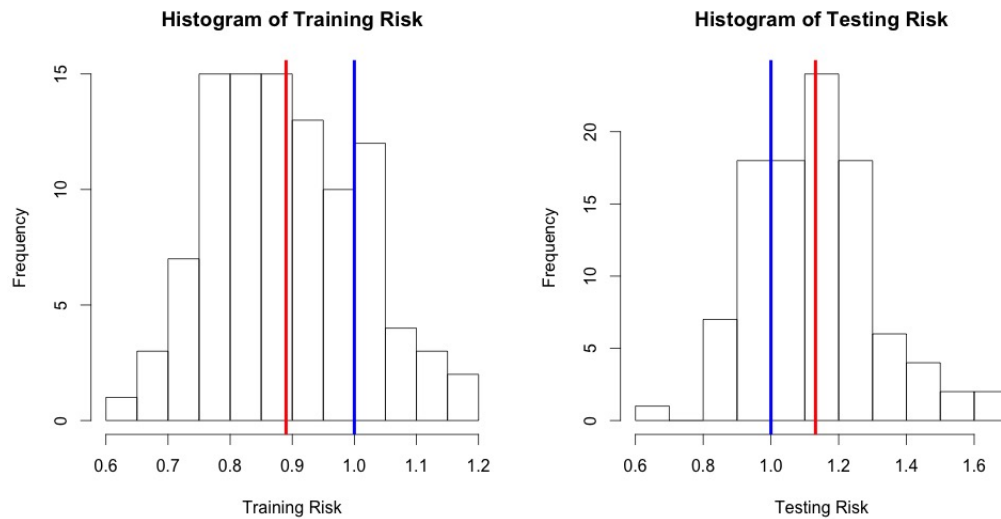
1 generateData = function(n, m, p){
2   #Takes in training and testing sample sizes and number of predictors.
3   #Generates simple noise data
4   yTrain = rnorm(n) #Responses
5   xTrain = matrix(rnorm(n*p), ncol=p) #Predictors
6
7   yTest = rnorm(m) #Responses
8   xTest = matrix(rnorm(m*p), ncol=p) #Predictors
9   list(yTrain=yTrain, xTrain=xTrain, yTest=yTest, xTest=xTest)
10 }
11
12 calcRisks = function(trainTest){
13   # Takes in training and testing data in a list.
14   #Unrolls into matrices, then calculates
15   # the training and testing errors (scalars). Returns a 2-vector.
16   yTrain = trainTest$yTrain
17   xTrain = trainTest$xTrain

```

```

18 yTest = trainTest$yTest
19 xTest = trainTest$xTest
20
21 #Compute training OLS estimate, get residuals/coeffs from it.
22 trainLM = lm(yTrain ~ xTrain)
23 trainResid = trainLM$resid
24 trainCoeffs = trainLM$coefficients
25
26 #Use training OLS estimates to make testing fits. Compute errors.
27 testFits = cbind(rep(1,nrow(xTest)), xTest) %*% trainCoeffs
28 testResid = yTest - testFits
29
30 riskTrain = 1/length(yTrain) * sum(trainResid^2)
31 riskTest = 1/length(yTest) * sum(testResid^2)
32 c(riskTrain, riskTest)
33 }
34
35 n = 100 #Train sample size
36 m = 100 #Test sample size
37 p = 10 #Number of predictors
38 nsim = 100 #Number of simulations
39
40 simErrors = t(sapply(1:nsim, function(i) calcRisks(generateData(n, m, p))))
41 par(mfrow=1:2)
42 hist(simErrors[,1], main='Histogram of Training Risk', xlab='Training Risk')
43 abline(v = mean(simErrors[,1]), col='red', lwd=4)
44 abline(v = 1, col='blue', lwd=4)
45 hist(simErrors[,2], main='Histogram of Testing Risk', xlab='Testing Risk')
46 abline(v = mean(simErrors[,2]), col='red', lwd=4)
47 abline(v = 1, col='blue', lwd=4)


```



## Problem 2

- (a)
- (b)

### Problem 3




- (a) A possible rule would be to classify observations to  $Y = 1$  if  $P(Y = 1) = \text{sigmoid}(\hat{\beta}^T X) \geq 0.5$ . This is equivalent to classifying to  $Y = 1$  if  $\hat{\beta}^T X \geq 0$ . Using our logistic regression, classify to  $Y = 1$  if  $0.127 + 0.063 \cdot x_1 + 1.390 \cdot x_2 \geq 0$ . The confusion matrix:

	Predict 0	Predict 1
Actual 0	308	87
Actual 1	101	304

The misclassification rate is  $(101 + 87)/800 = .235$ .

- (b) Classification rule drawn on the plot, corresponds to the line  $x_2 = -0.063x_1/1.390 - 0.127/1.390$ .



(c)

(d)

(e)

(f)

(g)